

Deep Learning for Visual Computing

Marcin Konefal, Daria Lazic

January 10, 2019

1 Assignment 2

1.1 Part 1: Gradient Descent

The gradient descent algorithm in general in deep learning is used to minimize a loss function that is non-linear in its parameters, like the cross entropy loss (very popular in classification). In the stochastic gradient descent, by definition parameters of a loss function are updated using one sample. In practice however, the batch size is usually 64, 128 or 256. In the assignment three different functions were tested. In an iterative process function point coordinates (which correspond to the parameters of the function) were updated using the gradient at the current position (deduced from calculated Hessian) multiplied with the learning rate. To evaluate the function at the current location, bilinear interpolation was used for approximation of non-integer coordinates. The algorithm has the most interesting behaviour for the Schaffer function, which has multiple sharp local minima. For other functions, the algorithm monotonously converges to the global minimum (at least if it is possible from the passed starting points). For the schaffer function however, the algorithm struggles to escape the local minimum, it is also much more sensitive to parameter change. For all three functions a learning rate of 10 works best. However, increasing the learning rate speeds up the optimization but can lead to overshooting. The same is true for epsilon, which is used for computing the numeric gradient descent. B builds up momentum if successive gradients are similar and thus improves the speed of convergence. A β of 0.9 turned out to be the most appropriate trade-off between speed and convergence. Nesterov momentum, together with a momentum parameter β of 0.9 gave best results. The stopping condition defines the value of the gradient at which the algorithm is stopped and depending on the function should be close to zero.

1.2 Part 2: Convolutional Neural Networks

The following newtork architecture was used:

```
class CNN(nn.Module):

    def __init__(self):
        super(CNN, self).__init__()

        # Layer 1, images are 32x32
        # Input channels = 3, output channels = 32
        # Padding for stride = 1 is calculated by (K-1)/2
        self._cn1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        # Layer 2, after pooling images are 16x16
        self._cn2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        # Layer 3, after pooling images are 8x8
        self._cn3 = nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1)
        # Layer 4, after pooling images are 4x4 - with padding 6x6
        self._cn4 = nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1)
        # Layer 4, after pooling images are 2x2 - with padding 4x4
        self._cn5 = nn.Conv2d(256, 512, kernel_size=3, stride=1, padding=1)

        self._pool = nn.MaxPool2d(kernel_size=2, stride=2, padding=0)
        self._global_avg_pool = nn.AdaptiveAvgPool2d(1) # 2 is the kernel size

        # 512*1*1 input features, 2 output features
        self._fc1 = nn.Linear(512, 2)

    def forward(self, data):
        data = F.relu(self._cn1(data))
        data = self._pool(data)
        data = F.relu(self._cn2(data))
        data = self._pool(data)
        data = F.relu(self._cn3(data))
        data = self._pool(data)
        data = F.relu(self._cn4(data))
        data = self._pool(data)
        data = F.relu(self._cn5(data))
        # global average
        data = self._global_avg_pool(data)
        data = data.view(-1, 512*1*1)
        data = self._fc1(data)

    return data
```

Listing 1: CNN

As can be seen in Listing 1, we used the basic design recipe presented in the lecture, which gave the best results. The first layer is a convolutional layer with a kernel size of 3 and stride 1. Padding for stride = 1 can be calculated by

$$padding = (kernelsize - 1)/2 \quad (1)$$

By using stride of 1 and padding of 1, the size of the image stays the same after convolution. Stride=2 would be too aggressive. Since our images have a size of WxH of 32x32 a kernel size of 3 makes sense (makes it possible to go deeper). After every conv layer, a ReLU activation function (sets all negative values to zero) and a pooling layer follow. Pooling is done by 2x2 max-pooling and stride of 2, which returns the maximum value in a window of 2x2 pixels. Width and height are reduced by a factor of 2 after each pooling layer. In the first convolutional layer we have three input channels (RGB) and 32 output channels, which correspond to 32 feature maps. In each subsequent block the number of feature maps is increased by a factor of 2. Since width and height of the image are reduced by a factor of 2 after each block, we used 4 subsequent blocks. After the last max-pooling layer, samples have a size of WxHxD of 2x2x256. After padding, WxH is 4x4 and thus a last conv layer with a kernel size of 3 and 512 output channels can be used (output

samples are of size $2 \times 2 \times 512$). The last pooling layer applies a global average pooling, which returns an average value for each feature map. So samples that are fed into the linear layer are of size $1 \times 1 \times 512$. The linear layer has as many neurons as classes (2 in our case).

As a loss function, cross entropy loss was used. The applied optimizer is an SGD with nesterov momentum and β of 0.9. Weight decay can be passed to the algorithm as an argument. The sequence of the training starts with a forward pass (Listing 1), followed by a backward pass to compute the gradients and ends with an optimization step using SGD to update model parameters (feature maps). A softmax layer is applied to the output of the linear layer for predicting probability class scores.




Figure 1: Training CNN without data augmentation.

Using a batch size of 128 samples and using input normalization, a learning rate of 0.1 turned out to be the best. The training loss goes to zero very quickly, while the validation accuracy reaches a maximum of 0.89 (Fig. 1).

1.3 Data Augmentation and Regularization

Data augmentation and regularization are used to reduce overfitting and 'explosion' of model parameters. While training loss decreases steadily with the number of epochs, validation loss never reaches the performance of the training loss and starts to increase at some point, because the model is overfitted to training data. This is because of the fact, that training data is limited and thus the model does not generalize well to unseen data. In addition to data augmentation and regularization, early stopping can be used, which saves the current state of the parameters if the highest validation performance is achieved (Listing 2).

```
if v_accuracy > best_accuracy:
    torch.save(cnn_net.state_dict(), os.path.join(os.getcwd(), 'best_model.pth'))
    best_accuracy = v_accuracy
```

Listing 2: Early stopping

In our assignment, we applied random cropping and flipping (Listing 3) for data augmentation, but also other transformations can be used like random scaling.

```
def hflip() -> Op:
    def op(sample: np.ndarray) -> np.ndarray:
        if random.random() < 0.5:
            return np.fliplr(sample)
        else:
            return sample
    return op

def rcrop(sz: int, pad: int, pad_mode: str) -> Op:
    def op(sample: np.ndarray) -> np.ndarray:
        if pad > 0:
            img = np.pad(sample, ((pad, pad), (pad, pad), (0, 0)), pad_mode)
        else:
            img = sample

        origin_size = img.shape
        cropped_size = (origin_size[0] - sz, origin_size[1] - sz)
        crop_point = (random.randrange(cropped_size[0]), random.randrange(
            cropped_size[1]))

        cropped_img = img[crop_point[0]: crop_point[0] + sz, crop_point[1]:
            crop_point[1] + sz]

        if cropped_img.shape[0] > origin_size[0] or cropped_img.shape[1] >
            origin_size[1]:
            raise ValueError("Image after cropping and padding has size: " +
                str(cropped_img.shape) +
                ", its size is bigger then size of origin picture: "
                + str(origin_size) + ".")
        return cropped_img
    return op
```

Listing 3: Flipping and Cropping

Regularization improves the validation and test performance with the expense of training performance, because of the possible decrease in model variance/ flexibility. In the assignment we used weight decay-which is incorporated in the optimizer-for regularization (penalizes large weights), but there are also other methods for regularization like dropout (output zero with probability p).

We tested different methods and combination of methods to reduce overfitting (Fig. 2, Fig. 3).




Figure 2: Training loss using different data augmentation and regularization methods. Without DA - without data augmentation and regularization- corresponds to part 2; Flipping - flipping only; Cropping - cropping only; Flipping + Cropping - flipping and cropping only; WD 10^{-4} - weight decay = 10^{-4} ; WD 10^{-5} - weight decay = 10^{-5} ; WD 10^{-6} - weight decay = 10^{-6} ; DA + WD: 10^{-4} - weight decay = 10^{-4} with flipping and cropping; DA + WD: 10^{-5} - weight decay = 10^{-5} with flipping and cropping.




Figure 3: Validation accuracy using different data augmentation and regularization methods. Without DA - without data augmentation and regularization- corresponds to part 2; Flipping - flipping only; Cropping - cropping only; Flipping + Cropping - flipping and cropping only; WD 10^{-4} - weight decay = 10^{-4} ; WD 10^{-5} - weight decay = 10^{-5} ; WD 10^{-6} - weight decay = 10^{-6} ; DA + WD: 10^{-4} - weight decay = 10^{-4} with flipping and cropping; DA + WD: 10^{-5} - weight decay = 10^{-5} with flipping and cropping.

The training loss goes to zero fast for all methods. With weight decay of 10^{-5} and 10^{-6} it stays at zero and there are less fluctuations. Cropping only has the worst performance concerning training loss. Validation accuracy reaches the highest value using data augmentation (flipping, cropping) and a weight decay of 10^{-5} . This maximum value of 0.921875 is reached already after 12 epochs. This model was thus considered the best model and model parameters were saved (Listing 2). The worse performance of using the model without any kind of augmentation or regularization is very prominent for validation accuracy.

1.4 Transfer Learning

For transfer learning we used pretrained torchvision models trained on the ImageNet data set. Instead of finetuning (retraining the whole model), we used feature extraction to only update the final layer weights. Therefore, we first initialized the model and then reshaped the final layer to have the same number of outputs as the number of classes in the dataset. There are two things that have to be considered. First of all, we have to define for the optimizer which parameters we want to update during training. This can be done by setting `param.requires_grad` to `True` for the layers that need to be updated.

```
def set_parameter_requires_grad(model, feature_extracting):
    if feature_extracting:
        for param in model.parameters():
            param.requires_grad = False
```

Listing 4: Selection of parameters which need to be updated.

In our case we set `requires_grad` for all parameters to false and then initialize a new last layer and by default the new parameters have `.requires_grad=True`. So only the new layer's parameters will be updated.

```
def initialize_model(model_name: str, num_classes: int, feature_extract: bool,
                    use_pretrained: bool=True):
    model_ft = None
    input_size = 0
    if model_name == "cnn":
        model_ft = CNN()
        set_parameter_requires_grad(model_ft, False)
        num_fts = model_ft._fc1.in_features
        model_ft._fc1 = nn.Linear(num_fts, num_classes)
        input_size = 32

    elif model_name == "resnet":
        """ Resnet18 """
        model_ft = models.resnet18(pretrained=use_pretrained)
        set_parameter_requires_grad(model_ft, feature_extract)
        num_fts = model_ft.fc.in_features
        model_ft.fc = nn.Linear(num_fts, num_classes)
        input_size = 224

    .....
    ....
    ..
```

Listing 5: Model initialization

The second thing that has to be considered is that inception_v3 model requires the input size to be (299, 299), whereas the other torchvision models expect an input size of (224, 224). Thus we added a resize function to ops.py.

```
def resize(val: int) -> Op:
    def op(sample: np.ndarray) -> np.ndarray:
        return ski_transform.resize(sample, (3, val, val), mode='reflect')

    return op
```

Listing 6: Resizing input dimensions

We also added an is_inception flag to the CNNClassifier, which is set to False by default. This flag is used because the Inception V3 model, uses an auxiliary and a final output, which are both incorporated into the model loss.

```
if self._is_inception:
    outputs, aux_outputs = self._net(train_data)
    loss1 = self._loss(outputs, train_labels)
    loss2 = self._loss(aux_outputs, train_labels)
    loss_train = loss1 + 0.4 * loss2
```

Listing 7: Auxiliary and final output in case of inception v3 model

We implemented the following models in the code: ResNet, AlexNet, VGG11_bn, SqueezeNet, DenseNet, Inception v3. Because the server was loaded and we were limited by time, we didn't try out any data augmentation or weight decay and only trained the ResNet and AlexNet for 10 epochs. We chose a value of 0.001 for the learning rate.




Figure 4: Training loss compared using different models for 10 epochs. CNN- our CNN model as in part 1; AlexNet/ResNet- torchvision model AlexNet/ResNet without any data augmentation or regularization.




Figure 5: Validation accuracy compared using different models for 10 epochs. CNN- our CNN model as in part 1; AlexNet/ResNet- torchvision model AlexNet/ResNet without any data augmentation or regularization.

The training loss with AlexNet or ResNet does not decrease as fast as with the CNN trained on our Cifar10 images but stops sinking at around 0.4. This might be due to the low number of epochs or due to the resizing of images from 32x32 to 224x224/299x299 or due to the kind of normalization we use. ResNet performs better than AlexNet concerning training loss and is much better than our CNN model concerning validation accuracy in the first 10 epochs. AlexNet has higher validation accuracy than CNN in the beginning but levels out after a few epochs. The training has to be done for more epochs to draw a meaningful conclusion.

1.5 Training with adapted accuracy update

Since we had all of the trainings already finished using the old version of test.py and received the correction of the first assignment afterwards, we agreed in the lecture to only do one more training with the modified test.py file. The wrong accuracy update overwrote the number of correct predictions and the overall number of predictions instead of summing it up. A new class object of accuracy is now initialized with each epoch and accuracy is updated (using correct update version) on the same validation batch (using 128 samples) after each training batch instead of resetting accuracy after each epoch and only calculating accuracy after the last training batch of the epoch. Instead of one more training, we did 4 in the end because we wanted to know what the performance of the models is using the adapted accuracy evaluation. Best validation accuracy of each trained model (Fig. 6, 7).

- CNN - no data augmenation - no regularization - learning rate 0.1 : 0.8035
- CNN - random flipping and cropping and WD of 0.00001 and learning rate of 0.1: 0.8644
- ResNet - no data augmentation - no regularization - learning rate 0.001 : 0.8589
- ResNet - random flipping and cropping and WD of 0.00001 and learning rate of 0.001: 0.86321

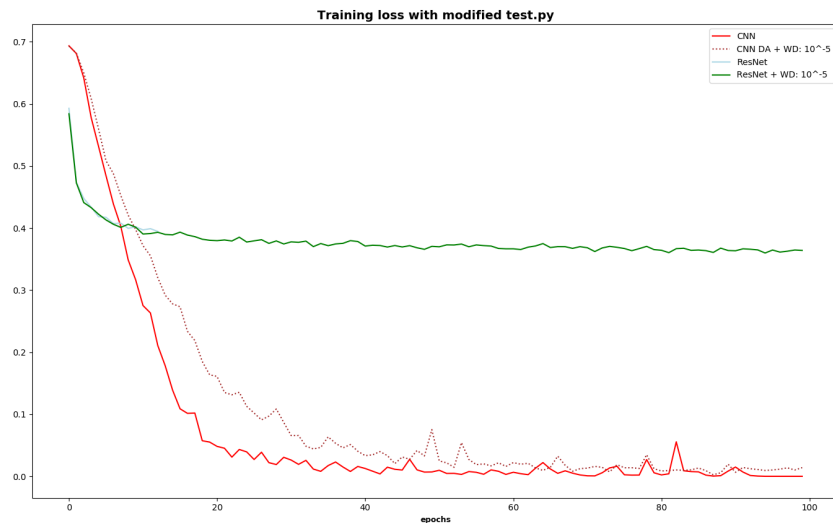


Figure 6: Training loss using different data augmentation and regularization methods. CNN - without data augmentation and regularization; CNN DA + WD: 10^{-5} - CNN with weight decay = 10^{-5} with flipping and cropping, ResNet - ResNet without data augmentation and regularization; ResNet DA + WD: 10^{-5} - ResNet with weight decay = 10^{-5} with flipping and cropping




Figure 7: Validation accuracy using different data augmentation and regularization methods. CNN - without data augmentation and regularization; CNN DA + WD: 10^{-5} - CNN with weight decay = 10^{-5} with flipping and cropping, ResNet - ResNet without data augmentation and regularization; ResNet DA + WD: 10^{-5} - ResNet with weight decay = 10^{-5} with flipping and cropping

It can be clearly seen, that validation accuracy is much better using data augmentation and regularization. Resnet does not outperform our CNN even with data augmentation and regularization - here training was stopped after 12 epochs, because validation accuracy was starting to decrease.